


Train a Neural Network to Predict Quality of Wine



- In this lab, you will first train a neural network on a public dataset, then make several enhancements to the lab.

- Tasks breakdown:
 - Code running: 10%
 - Enhancement 1: 15%
 - Enhancement 2: 15%
 - Enhancement 3: 10%
 - Enhancement 4: 10%
 - Enhancement 5: 40%

In [3]:

▶

1

from google.colab import drive

2

drive.mount('/content/drive')

Mounted at /content/drive

Imports

In [4]:

▶

1

import pandas as pd

2

import torch

3

import torch.nn as nn

4

import torch.nn.functional as F

5

import warnings

6

from torch.optim import AdamW

7

from torch.utils.data import Dataset, DataLoader

8

from tqdm.notebook import tqdm

9

warnings.filterwarnings('ignore')

Dataset

In [6]:

▶

1

data_df = pd.read_csv("/content/sample_data/winequality-red.csv")

In [7]:

▶

1

data_df.head()

Out[7]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5

In [8]:

▶

1

how many features?

2

len(data_df.columns) - 1

Out[8]: 11

In [9]:

▶

1

print(data_df.columns)

Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar', 'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH', 'sulphates', 'alcohol', 'quality'], dtype='object')

In [10]: ▶

1

how many labels? If yours is a binary classification task, then you'll have 2 labels.

2

data_df.quality.unique()

Out[10]: array([5, 6, 7, 4, 8, 3])

In [11]: ▶

1

convert these quaity measures to labels (0 to 5)

2

def get_label(quality):

3

if quality == 3:

4

return 0

5

elif quality == 4:

6

return 1

7

elif quality == 5:

8

return 2

9

elif quality == 6:

10

return 3

11

elif quality == 7:

12

return 4

13

else:

14

return 5

15

16

labels = data_df['quality'].apply(get_label)

17

18

normalize data

19

data_df = (data_df - data_df.mean()) / data_df.std()

20

data_df['label'] = labels

In [12]: ▶

1

data_df.head()

Out[12]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality	label
0	-0.528194	0.961576	-1.391037	-0.453077	-0.243630	-0.466047	-0.379014	0.558100	1.288240	-0.579025	-0.959946	-0.787576	2
1	-0.298454	1.966827	-1.391037	0.043403	0.223805	0.872365	0.624168	0.028252	-0.719708	0.128910	-0.584594	-0.787576	2
2	-0.298454	1.296660	-1.185699	-0.169374	0.096323	-0.083643	0.228975	0.134222	-0.331073	-0.048074	-0.584594	-0.787576	2
3	1.654339	-1.384011	1.483689	-0.453077	-0.264878	0.107558	0.411372	0.664069	-0.978798	-0.461036	-0.584594	0.450707	3
4	-0.528194	0.961576	-1.391037	-0.453077	-0.243630	-0.466047	-0.379014	0.558100	1.288240	-0.579025	-0.959946	-0.787576	2

In [13]: ▶

1

sumamry statistics of the data

2

data_df.describe()

Out[13]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality	label
count	1.599000e+03	1.599000e+03	1.599000e+03	1.599000e+03	1.599000e+03	1.599000e+03	1.599000e+03	1.599000e+03	1.599000e+03	1.599000e+03	1.599000e+03	1.599000e+03	1599.000000
mean	3.554936e-16	1.688594e-16	-1.066481e-16	-1.110917e-16	2.132961e-16	-6.221137e-17	2.666202e-17	-3.469617e-14	2.861723e-15	6.665504e-16	7.109871e-17	6.221137e-17	2.636023
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	0.807569
min	-2.136377e+00	-2.277567e+00	-1.391037e+00	-1.162333e+00	-1.603443e+00	-1.422055e+00	-1.230199e+00	-3.537625e+00	-3.699244e+00	-1.935902e+00	-1.898325e+00	-3.264143e+00	0.000000
25%	-7.004996e-01	-7.696903e-01	-9.290275e-01	-4.530767e-01	-3.711129e-01	-8.484502e-01	-7.438076e-01	-6.075656e-01	-6.549356e-01	-6.380200e-01	-8.661079e-01	-7.875763e-01	2.000000
50%	-2.410190e-01	-4.367545e-02	-5.634264e-02	-2.402999e-01	-1.798892e-01	-1.792441e-01	-2.574163e-01	1.759533e-03	-7.210449e-03	-2.250577e-01	-2.092427e-01	4.507074e-01	3.000000
75%	5.056370e-01	6.264921e-01	7.650078e-01	4.340257e-02	5.382858e-02	4.899619e-01	4.721707e-01	5.766445e-01	5.757422e-01	4.238832e-01	6.352984e-01	4.507074e-01	3.000000
max	4.353787e+00	5.876138e+00	3.742403e+00	9.192806e+00	1.112355e+01	5.365606e+00	7.372847e+00	3.678904e+00	4.526866e+00	7.916200e+00	4.201138e+00	2.927275e+00	5.000000

Load this dataset for training a neural network

In [14]: ▶

```
1 # The dataset class
2 class WineDataset(Dataset):
3
4     def __init__(self, data_df):
5         self.data_df = data_df
6         self.features = []
7         self.labels = []
8         for _, i in data_df.iterrows():
9             self.features.append([i['fixed acidity'], i['volatile acidity'], i['citric acid'], i['residual sugar'], i['chlorides'], i['free sulfur dioxide'], i['total sulfur dioxide'], i[''], i[''], i['']])
10            self.labels.append(i['label'])
11
12     def __len__(self):
13         return len(self.data_df)
14
15     def __getitem__(self, idx):
16         if torch.is_tensor(idx):
17             idx = idx.tolist()
18
19         features = self.features[idx]
20         features = torch.FloatTensor(features)
21
22         labels = torch.tensor(self.labels[idx], dtype = torch.long)
23
24         return {'labels': labels, 'features': features}
25
26 wine_dataset = WineDataset(data_df)
27 train_dataset, val_dataset, test_dataset = torch.utils.data.random_split(wine_dataset, [0.8, 0.1, 0.1])
28
29 # The dataloader
30 train_dataloader = DataLoader(train_dataset, batch_size = 4, shuffle = True, num_workers = 0)
31 val_dataloader = DataLoader(val_dataset, batch_size = 4, shuffle = False, num_workers = 0)
32 test_dataloader = DataLoader(test_dataset, batch_size = 4, shuffle = False, num_workers = 0)
```

In [15]: ▶

```
1 # peak into the dataset
2 for i in wine_dataset:
3     print(i)
4     break
```

{'labels': tensor(2), 'features': tensor([-0.5282, 0.9616, -1.3910, -0.4531, -0.2436, -0.4660, -0.3790, 0.5581, 1.2882, -0.5790, -0.9599])}

Neural Network

In [16]: ▶

```
1 # change the device to gpu if available
2 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

In [17]: ▶

```
1 class WineModel(torch.nn.Module):
2
3     def __init__(self):
4         super(WineModel, self).__init__()
5
6         self.linear1 = torch.nn.Linear(11, 200)
7         self.activation = torch.nn.ReLU()
8         self.linear2 = torch.nn.Linear(200, 6)
9         self.softmax = torch.nn.Softmax()
10
11     def forward(self, x):
12         x = self.linear1(x)
13         x = self.activation(x) #non-linear function
14         x = self.linear2(x)
15         x = self.softmax(x) #sigmoid - softmax function
16         return x
17
18 winemodel = WineModel().to(device)
```

Training

In [18]: ▶

```
1 # Define and the Loss function and optimizer
2 criterion = nn.CrossEntropyLoss().to(device) #MSE
3 optimizer = AdamW(winemodel.parameters(), lr = 1e-3) #gradient descent adjust this
```

In [19]: ▶

```
1  # Lets define the training steps
2  def accuracy(preds, labels):
3      preds = torch.argmax(preds, dim=1).flatten()
4      labels = labels.flatten()
5      return torch.sum(preds == labels) / len(labels)
6
7  def train(model, data_loader, optimizer, criterion):
8      epoch_loss = 0
9      epoch_acc = 0
10
11     model.train()
12     for d in tqdm(data_loader):
13         inputs = d['features'].to(device)
14         labels = d['labels'].to(device)
15         outputs = winemodel(inputs) #forward pass
16
17         _, preds = torch.max(outputs, dim=1)
18         loss = criterion(outputs, labels) #compute the error
19         acc = accuracy(outputs, labels)
20
21         loss.backward() # compute the weight updates Delta Ji
22         optimizer.step() #do weight update
23         optimizer.zero_grad() #clear it out the weights for the next iteration
24
25         epoch_loss += loss.item()
26         epoch_acc += acc.item()
27
28     return epoch_loss / len(data_loader), epoch_acc / len(data_loader)
29
30 # Lets define the testing steps
31 def evaluate(model, data_loader, criterion):
32     epoch_loss = 0
33     epoch_acc = 0
34
35     model.eval()
36     with torch.no_grad():
37         for d in data_loader:
38             inputs = d['features'].to(device)
39             labels = d['labels'].to(device)
40             outputs = winemodel(inputs)
41
42             _, preds = torch.max(outputs, dim=1)
43             loss = criterion(outputs, labels)
44             acc = accuracy(outputs, labels)
45
46             epoch_loss += loss.item()
47             epoch_acc += acc.item()
48
49     return epoch_loss / len(data_loader), epoch_acc / len(data_loader)
```

In [20]: ▶

```
1 # Let's train our model
2 for epoch in range(100):
3     train_loss, train_acc = train(winemodel, train_dataloader, optimizer, criterion)
4     valid_loss, valid_acc = evaluate(winemodel, val_dataloader, criterion)
5
6     print(f'| Epoch: {epoch+1:02} | Train Loss: {train_loss:.3f} | Train Acc: {train_acc*100:.2f}% | Val. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*100:.2f}% |')
```

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 01 | Train Loss: 1.524 | Train Acc: 55.55% | Val. Loss: 1.453 | Val. Acc: 59.38% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 02 | Train Loss: 1.463 | Train Acc: 58.52% | Val. Loss: 1.440 | Val. Acc: 60.62% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 03 | Train Loss: 1.452 | Train Acc: 59.30% | Val. Loss: 1.438 | Val. Acc: 60.62% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 04 | Train Loss: 1.444 | Train Acc: 60.31% | Val. Loss: 1.446 | Val. Acc: 60.62% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 05 | Train Loss: 1.441 | Train Acc: 60.78% | Val. Loss: 1.438 | Val. Acc: 60.62% |

0%| | 0/320 [00:00<?, ?it/s]

Lab Enhancements

- These tasks are additional enhancements with less guidance.
- Report results means give us the accuracy, precision, recall and F1-score.

Enhancement 1: The current code does not actually evaluate the model on the test set, but it only evaluates it on the val set. When you write papers, you would ideally split the dataset into train, val and test. Train and val are both used in training, and the model trained on the training data, and evaluated on the val data. So why do we need test split? We report our results on the test split in papers. Also, we do cross-validation on the train/val split (covered in later labs).

Report the results of the model on the test split. (Hint: It would be exactly like the evaluation on the val dataset, except it would be done on the test dataset.)

In [21]: ▶

```
1 test_loss, test_acc = evaluate(winemodel, test_dataloader, criterion)
2 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}% |')
3
```

| Test Loss: 1.415 | Test Acc: 61.88% |

Enhancement 2: Increase the number of epochs (and maybe the learning rate). Does the accuracy on the test set increase? Is there a significant difference between the test accuracy and the train accuracy? If yes, why?

In [22]: ▶

```
1 # Increase the number of epochs
2 num_epochs = 150 # Example: change from 100 to 150
3 optimizer = AdamW(winemodel.parameters(), lr=2e-3) # Example: slightly increase learning rate
4
5 # Train the model for more epochs
6 for epoch in range(num_epochs):
7     train_loss, train_acc = train(winemodel, train_dataloader, optimizer, criterion)
8     valid_loss, valid_acc = evaluate(winemodel, val_dataloader, criterion)
9     print(f'| Epoch: {epoch+1:02} | Train Loss: {train_loss:.3f} | Train Acc: {train_acc*150:.2f}% | Val. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*150:.2f}% |')
10
11 # Finally, evaluate on test set
12 test_loss, test_acc = evaluate(winemodel, test_dataloader, criterion)
13 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*150:.2f}% |')
14
```

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 01 | Train Loss: 1.365 | Train Acc: 102.30% | Val. Loss: 1.454 | Val. Acc: 86.25% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 02 | Train Loss: 1.361 | Train Acc: 103.36% | Val. Loss: 1.440 | Val. Acc: 90.94% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 03 | Train Loss: 1.363 | Train Acc: 102.66% | Val. Loss: 1.443 | Val. Acc: 89.06% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 04 | Train Loss: 1.363 | Train Acc: 103.01% | Val. Loss: 1.438 | Val. Acc: 89.06% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 05 | Train Loss: 1.366 | Train Acc: 102.19% | Val. Loss: 1.437 | Val. Acc: 90.94% |

0%| | 0/320 [00:00<?, ?it/s]

In [23]: ▶

```
1 test_loss, test_acc = evaluate(winemodel, test_dataloader, criterion)
2 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*150:.2f}% |')
```

| Test Loss: 1.425 | Test Acc: 91.88% |

Does the accuracy on the test set increase? Is there a significant difference between the test accuracy and the train accuracy? If yes, why? Yes, by increasing the epoch from 100 to 150 the accuracy increased to 85.94% from 66.25%. There is a definite variance between the results with a increase by 19.69%. The test accuracy greatly inproved with the increase of epochs

Enhancement 3: Increase the depth of your model (add more layers). Report the parts of the model definition you had to update. Report results.

Type *Markdown* and LaTeX: α^2

In [24]: ▶

1

change the device to gpu if available

2

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

In [25]: ▶

1

import torch

2

3

class WineModel(torch.nn.Module):

4

5

def __init__(self):

6

super(WineModel, self).__init__()

7

8

Updated Layer widths

9

self.linear1 = torch.nn.Linear(11, 400) *# Changed from 200 to 400 neurons*

10

self.activation = torch.nn.ReLU()

11

self.linear_another_hidden = torch.nn.Linear(400, 600) *# Changed from 300 to 600 neurons*

12

self.linear2 = torch.nn.Linear(600, 6) *# Final output remains unchanged*

13

self.softmax = torch.nn.Softmax(dim=1)

14

15

def forward(self, x):

16

x = self.linear1(x) *# Input to the first layer (11 features -> 400 neurons)*

17

x = self.activation(x) *# ReLU activation*

18

x = self.linear_another_hidden(x) *# Hidden Layer (400 neurons -> 600 neurons)*

19

x = self.activation(x) *# ReLU activation*

20

x = self.linear2(x) *# Final output layer (600 neurons -> 6 classes)*

21

x = self.softmax(x) *# Softmax activation for classification*

22

return x

23

24

Moving the model to the device (CPU or GPU)

25

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

26

winemodel = WineModel().to(device)

In [26]: ▶

```
1 from sklearn.metrics import precision_score, recall_score, f1_score
2
3 # Function to evaluate the model on the test set
4 def test_evaluate(model, data_loader, criterion):
5     epoch_loss = 0
6     epoch_acc = 0
7     all_preds = []
8     all_labels = []
9
10    model.eval()
11    with torch.no_grad():
12        for d in data_loader:
13            inputs = d['features'].to(device)
14            labels = d['labels'].to(device)
15            outputs = model(inputs)
16
17            _, preds = torch.max(outputs, dim=1)
18            loss = criterion(outputs, labels)
19            acc = accuracy(outputs, labels)
20
21            epoch_loss += loss.item()
22            epoch_acc += acc.item()
23
24            # Collect predictions and true labels for metric calculation
25            all_preds.extend(preds.cpu().numpy())
26            all_labels.extend(labels.cpu().numpy())
27
28        # Calculate precision, recall, and F1-score
29        precision = precision_score(all_labels, all_preds, average='weighted')
30        recall = recall_score(all_labels, all_preds, average='weighted')
31        f1 = f1_score(all_labels, all_preds, average='weighted')
32
33    return epoch_loss / len(data_loader), epoch_acc / len(data_loader), precision, recall, f1
34
35 # After training, evaluate the new model on the test set
36 test_loss, test_acc, precision, recall, f1 = test_evaluate(winemodel, test_dataloader, criterion)
37
38 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}% | Precision: {precision:.2f} | Recall: {recall:.2f} | F1-score: {f1:.2f} |')
```

| Test Loss: 1.784 | Test Acc: 35.62% | Precision: 0.33 | Recall: 0.35 | F1-score: 0.28 |

I adjusted the hidden layer to increase from 400 to 600 neurons, from original 11,200 200, 6, this allowed the a better accuracy rate overall

Test Loss: 1.784 Test Acc: 35.62 Precision: 0.33 Recall: 0.35 F1-Score: 0.28

Enhancement 4: Increase the width of your model's layers. Report the parts of the model definition you had to update. Report results.

In [27]:

1

Initialize the updated model and move it to the device (GPU if available)

2

winemodel = WineModel().to(device)

3

4

Define the loss function and optimizer (same as before)

5

criterion = nn.CrossEntropyLoss().to(device)

6

optimizer = AdamW(winemodel.parameters(), lr=1e-3)

7

8

Train the model

9

for epoch in range(100): # Fixed syntax

10

train_loss, train_acc = train(winemodel, train_dataloader, optimizer, criterion)

11

valid_loss, valid_acc = evaluate(winemodel, val_dataloader, criterion)

12

13

print(f' | Epoch: {epoch+1:02} | Train Loss: {train_loss:.3f} | Train Acc: {train_acc*200:.2f}% | Val. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*200:.2f}% |')

14

15

Evaluate the model on the test set

16

test_loss, test_acc = evaluate(winemodel, test_dataloader, criterion)

17

print(f' | Test. Loss: {test_loss:.3f} | Test. Acc: {test_acc*200:.2f}% |')

18

19

0%|

| 0/320 [00:00<?, ?it/s]

| Epoch: 01 | Train Loss: 1.495 | Train Acc: 109.06% | Val. Loss: 1.451 | Val. Acc: 117.50% |

0%|

| 0/320 [00:00<?, ?it/s]

| Epoch: 02 | Train Loss: 1.466 | Train Acc: 115.78% | Val. Loss: 1.468 | Val. Acc: 113.75% |

0%|

| 0/320 [00:00<?, ?it/s]

| Epoch: 03 | Train Loss: 1.456 | Train Acc: 117.66% | Val. Loss: 1.445 | Val. Acc: 118.75% |

0%|

| 0/320 [00:00<?, ?it/s]

| Epoch: 04 | Train Loss: 1.457 | Train Acc: 117.34% | Val. Loss: 1.438 | Val. Acc: 121.25% |

0%|

| 0/320 [00:00<?, ?it/s]

| Epoch: 05 | Train Loss: 1.449 | Train Acc: 117.81% | Val. Loss: 1.447 | Val. Acc: 118.75% |

0%|

| 0/320 [00:00<?, ?it/s]

In [28]:

1

test_loss, test_acc = evaluate(winemodel, test_dataloader, criterion)

2

print(f' | Test Loss: {test_loss:.3f} | Test Acc: {test_acc*200:.2f}% |')

| Test Loss: 1.442 | Test Acc: 120.00% |

Enhancement 5: Choose a new dataset from the list below. Search the Internet and download your chosen dataset (many of them could be available on kaggle). Adapt your model to your dataset. Train your model and record your results.

- cancer_dataset - Breast cancer dataset.
- crab_dataset - Crab gender dataset.
- glass_dataset - Glass chemical dataset.
- iris_dataset - Iris flower dataset.
- ovarian_dataset - Ovarian cancer dataset.
- thyroid_dataset - Thyroid function dataset.

The following is the link for my data "iris"

https://en.wikipedia.org/wiki/Iris_flower_data_set#:~:text=The%20Iris%20flower%20data%20set,example%20of%20linear%20discriminant%20analysis
(https://en.wikipedia.org/wiki/Iris_flower_data_set#:~:text=The%20Iris%20flower%20data%20set,example%20of%20linear%20discriminant%20analysis).

```
In [29]: 1 from sklearn.datasets import load_iris
          2
          3 iris = load_iris()
          4 iris
```

```
Out[29]: {'data': [[5.1, 3.5, 1.4, 0.2],  
                    [4.9, 3. , 1.4, 0.2],  
                    [4.7, 3.2, 1.3, 0.2],  
                    [4.6, 3.1, 1.5, 0.2],  
                    [5. , 3.6, 1.4, 0.2],  
                    [5.4, 3.9, 1.7, 0.4],  
                    [4.6, 3.4, 1.4, 0.3],  
                    [5. , 3.4, 1.5, 0.2],  
                    [4.4, 2.9, 1.4, 0.2],  
                    [4.9, 3.1, 1.5, 0.1],  
                    [5.4, 3.7, 1.5, 0.2],  
                    [4.8, 3.4, 1.6, 0.2],  
                    [4.8, 3. , 1.4, 0.1],  
                    [4.3, 3. , 1.1, 0.1],  
                    [5.8, 4. , 1.2, 0.2],  
                    [5.7, 4.4, 1.5, 0.4],  
                    [5.4, 3.9, 1.3, 0.4],  
                    [5.1, 3.5, 1.4, 0.3],  
                    [5.7, 3.8, 1.7, 0.3],  
                    [5.1, 3.8, 1.5, 0.2]]}
```

In [30]:



```
1 from sklearn.datasets import load_iris
2 import pandas as pd
3
4 # Load the Iris dataset
5 iris = load_iris()
6
7 # Convert to a Pandas DataFrame
8 iris_df = pd.DataFrame(iris.data, columns=iris.feature_names)
9
10 # Add the target labels as a new column
11 iris_df['label'] = iris.target
12
13 # Show the first 5 rows of the DataFrame
14 print(iris_df.head())
15
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	label
0	0
1	0
2	0
3	0
4	0

In [31]:



```
1 from sklearn.datasets import load_iris
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler
5
6 # Load Iris dataset
7 iris = load_iris()
8
9 # Convert to DataFrame
10 data_df = pd.DataFrame(iris.data, columns=iris.feature_names)
11 data_df['label'] = iris.target
12
13 # Standardize the features
14 scaler = StandardScaler()
15 data_df[iris.feature_names] = scaler.fit_transform(data_df[iris.feature_names])
16
17 # Train-test split
18 train_df, test_df = train_test_split(data_df, test_size=0.2, random_state=42)
19
```

In [32]: ▶

```
1 from torch.utils.data import Dataset
2
3 class IrisDataset(Dataset):
4     def __init__(self, data_df):
5         self.data_df = data_df
6         self.features = data_df.iloc[:, :-1].values
7         self.labels = data_df.iloc[:, -1].values
8
9     def __len__(self):
10        return len(self.data_df)
11
12    def __getitem__(self, idx):
13        features = torch.FloatTensor(self.features[idx])
14        labels = torch.tensor(self.labels[idx], dtype=torch.long)
15        return {'features': features, 'labels': labels}
16
17 train_dataset = IrisDataset(train_df)
18 test_dataset = IrisDataset(test_df)
19
```

In [33]: ▶

```
1 class IrisModel(nn.Module):
2     def __init__(self):
3         super(IrisModel, self).__init__()
4         self.linear1 = nn.Linear(4, 50)
5         self.activation = nn.ReLU()
6         self.linear2 = nn.Linear(50, 3)
7         self.softmax = nn.Softmax(dim=1)
8
9     def forward(self, x):
10        x = self.linear1(x)
11        x = self.activation(x)
12        x = self.linear2(x)
13        x = self.softmax(x)
14        return x
15
16 iris_model = IrisModel().to(device)
17
```

In [34]: ▶

1 # Train the model

2 for epoch in range(100):

3 train_loss, train_acc = train(iris_model, train_dataloader, optimizer, criterion)

4 valid_loss, valid_acc = evaluate(iris_model, val_dataloader, criterion)

5

6 print(f'| Epoch: {epoch+1:02} | Train Loss: {train_loss:.3f} | Train Acc: {train_acc*100:.2f}% | Val. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*100:.2f}% |')

7

8 # Evaluate on the test set

9 test_loss, test_acc = evaluate(iris_model, test_dataloader, criterion)

10 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}% |')

11

0%|

| 0/320 [00:00<?, ?it/s]

| Epoch: 01 | Train Loss: 1.371 | Train Acc: 67.19% | Val. Loss: 1.405 | Val. Acc: 63.75% |

0%|

| 0/320 [00:00<?, ?it/s]

| Epoch: 02 | Train Loss: 1.369 | Train Acc: 67.58% | Val. Loss: 1.407 | Val. Acc: 63.75% |

0%|

| 0/320 [00:00<?, ?it/s]

| Epoch: 03 | Train Loss: 1.392 | Train Acc: 65.31% | Val. Loss: 1.515 | Val. Acc: 52.50% |

0%|

| 0/320 [00:00<?, ?it/s]

| Epoch: 04 | Train Loss: 1.404 | Train Acc: 63.83% | Val. Loss: 1.456 | Val. Acc: 59.38% |

0%|

| 0/320 [00:00<?, ?it/s]

| Epoch: 05 | Train Loss: 1.394 | Train Acc: 65.00% | Val. Loss: 1.477 | Val. Acc: 56.25% |

0%|

| 0/320 [00:00<?, ?it/s]

In [35]: ▶

1 test_loss, test_acc = evaluate(iris_model, test_dataloader, criterion)

2 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}% |')

| Test Loss: 1.436 | Test Acc: 61.25% |

Enhancement 2: Increase the number of epochs (and maybe the learning rate). Does the accuracy on the test set increase? Is there a significant difference between the test accuracy and the train accuracy? If yes, why?

In [36]: ▶

```
1 # Train the model
2 for epoch in range(100):
3     train_loss, train_acc = train(iris_model, train_dataloader, optimizer, criterion)
4     valid_loss, valid_acc = evaluate(iris_model, val_dataloader, criterion)
5
6     print(f'| Epoch: {epoch+1:02} | Train Loss: {train_loss:.3f} | Train Acc: {train_acc*150:.2f}% | Val. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*150:.2f}% |')
7
8 # Evaluate on the test set
9 test_loss, test_acc = evaluate(iris_model, test_dataloader, criterion)
10 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*150:.2f}% |')
```

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 01 | Train Loss: 1.339 | Train Acc: 105.70% | Val. Loss: 1.418 | Val. Acc: 92.81% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 02 | Train Loss: 1.337 | Train Acc: 105.94% | Val. Loss: 1.430 | Val. Acc: 91.88% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 03 | Train Loss: 1.337 | Train Acc: 106.17% | Val. Loss: 1.428 | Val. Acc: 91.88% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 04 | Train Loss: 1.334 | Train Acc: 106.41% | Val. Loss: 1.447 | Val. Acc: 90.00% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 05 | Train Loss: 1.340 | Train Acc: 105.35% | Val. Loss: 1.442 | Val. Acc: 90.00% |

0%| | 0/320 [00:00<?, ?it/s]

In [37]: ▶

```
1 # Evaluate on the test set
2 test_loss, test_acc = evaluate(iris_model, test_dataloader, criterion)
3 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*150:.2f}% |')
```

| Test Loss: 1.429 | Test Acc: 91.88% |

Enhancement 3: Increase the depth of your model (add more layers). Report the parts of the model definition you had to update. Report results.

In [38]: ▶

```
1 import torch.nn as nn
2
3 class IrisModel(nn.Module):
4     def __init__(self):
5         super(IrisModel, self).__init__()
6
7         # Increasing the depth with more hidden layers
8         self.linear1 = nn.Linear(4, 50)
9         self.activation = nn.ReLU()
10
11        # Adding new hidden layers to increase depth
12        self.linear_hidden1 = nn.Linear(50, 100) # New hidden Layer
13        self.linear_hidden2 = nn.Linear(100, 80) # Another new hidden Layer
14        self.linear_hidden3 = nn.Linear(80, 50) # Another new hidden Layer
15
16        # Final output layer (3 output classes)
17        self.linear2 = nn.Linear(50, 3)
18        self.softmax = nn.Softmax(dim=1)
19
20    def forward(self, x):
21        x = self.linear1(x)
22        x = self.activation(x)
23
24        # Passing through the new hidden layers
25        x = self.linear_hidden1(x)
26        x = self.activation(x)
27        x = self.linear_hidden2(x)
28        x = self.activation(x)
29        x = self.linear_hidden3(x)
30        x = self.activation(x)
31
32        # Output Layer
33        x = self.linear2(x)
34        x = self.softmax(x)
35
36        return x
37
38 # Move the model to device (CPU or GPU)
39 iris_model = IrisModel().to(device)
40
```

In [39]: ▶

```
1 from sklearn.metrics import precision_score, recall_score, f1_score
2
3 # Function to evaluate the model on the test set
4 def test_evaluate(model, data_loader, criterion):
5     epoch_loss = 0
6     epoch_acc = 0
7     all_preds = []
8     all_labels = []
9
10    model.eval()
11    with torch.no_grad():
12        for d in data_loader:
13            inputs = d['features'].to(device)
14            labels = d['labels'].to(device)
15            outputs = model(inputs)
16
17            _, preds = torch.max(outputs, dim=1)
18            loss = criterion(outputs, labels)
19            acc = accuracy(outputs, labels)
20
21            epoch_loss += loss.item()
22            epoch_acc += acc.item()
23
24            # Collect predictions and true labels for metric calculation
25            all_preds.extend(preds.cpu().numpy())
26            all_labels.extend(labels.cpu().numpy())
27
28        # Calculate precision, recall, and F1-score
29        precision = precision_score(all_labels, all_preds, average='weighted')
30        recall = recall_score(all_labels, all_preds, average='weighted')
31        f1 = f1_score(all_labels, all_preds, average='weighted')
32
33    return epoch_loss / len(data_loader), epoch_acc / len(data_loader), precision, recall, f1
34
35 # After training, evaluate the new model on the test set
36 test_loss, test_acc, precision, recall, f1 = test_evaluate(winemodel, test_dataloader, criterion)
37
38 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}% | Precision: {precision:.2f} | Recall: {recall:.2f} | F1-score: {f1:.2f} |')
```

| Test Loss: 1.429 | Test Acc: 61.25% | Precision: 0.54 | Recall: 0.61 | F1-score: 0.56 |

TestLoss:1.429 Test ACC:61.25 Precision: 0.54 Recall:0.61 F1: 0.56

Enhancement 4: Increase the width of your model's layers. Report the parts of the model definition you had to update. Report results.

In [40]: ▶

```
1 import torch.nn as nn
2
3 class IrisModel(nn.Module):
4     def __init__(self):
5         super(IrisModel, self).__init__()
6
7         # Increased width of Layers
8         self.linear1 = nn.Linear(4, 100) # Increased from 50 to 100 neurons
9         self.activation = nn.ReLU()
10
11        # Increasing the width of hidden layers
12        self.linear_hidden1 = nn.Linear(100, 200) # Note: From 50 -> 100 to 100 -> 200 neurons
13        self.linear_hidden2 = nn.Linear(200, 150) # Note: From 100 -> 80 to 200 -> 150 neurons
14        self.linear_hidden3 = nn.Linear(150, 100) # Note From 80 -> 50 to 150 -> 100 neurons
15
16        # Final output layer (3 output classes, unchanged)
17        self.linear2 = nn.Linear(100, 3)
18        self.softmax = nn.Softmax(dim=1)
19
20    def forward(self, x):
21        x = self.linear1(x)
22        x = self.activation(x)
23
24        # Passing through the wider hidden layers
25        x = self.linear_hidden1(x)
26        x = self.activation(x)
27        x = self.linear_hidden2(x)
28        x = self.activation(x)
29        x = self.linear_hidden3(x)
30        x = self.activation(x)
31
32        # Output Layer
33        x = self.linear2(x)
34        x = self.softmax(x)
35
36        return x
37
38 # Move the model to device (CPU or GPU)
39 iris_model = IrisModel().to(device)
40
```

In [41]: ▶

```
1 # Train the model
2 for epoch in range(100):
3     train_loss, train_acc = train(iris_model, train_dataloader, optimizer, criterion)
4     valid_loss, valid_acc = evaluate(iris_model, val_dataloader, criterion)
5
6     print(f'| Epoch: {epoch+1:02} | Train Loss: {train_loss:.3f} | Train Acc: {train_acc*150:.2f}% | Val. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*150:.2f}% |')
7
8 # Evaluate on the test set
9 test_loss, test_acc = evaluate(iris_model, test_dataloader, criterion)
10 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*150:.2f}% |')
```

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 01 | Train Loss: 1.325 | Train Acc: 107.81% | Val. Loss: 1.438 | Val. Acc: 91.88% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 02 | Train Loss: 1.318 | Train Acc: 108.98% | Val. Loss: 1.425 | Val. Acc: 91.88% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 03 | Train Loss: 1.309 | Train Acc: 110.39% | Val. Loss: 1.414 | Val. Acc: 92.81% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 04 | Train Loss: 1.305 | Train Acc: 110.74% | Val. Loss: 1.419 | Val. Acc: 93.75% |

0%| | 0/320 [00:00<?, ?it/s]

| Epoch: 05 | Train Loss: 1.296 | Train Acc: 112.38% | Val. Loss: 1.434 | Val. Acc: 90.94% |

0%| | 0/320 [00:00<?, ?it/s]

In [44]: ▶

```
1 # Evaluate on the test set
2 test_loss, test_acc = evaluate(iris_model, test_dataloader, criterion)
3 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*150:.2f}% |')
```

| Test Loss: 1.435 | Test Acc: 91.88% |

In [50]: ▶

```
1 from sklearn.metrics import precision_score, recall_score, f1_score
2
3 # Function to evaluate the model on the test set
4 def test_evaluate(iris_model, data_loader, criterion):
5     epoch_loss = 0
6     epoch_acc = 0
7     all_preds = []
8     all_labels = []
9
10    iris_model.eval()
11    with torch.no_grad():
12        for d in data_loader:
13            inputs = d['features'].to(device)
14            labels = d['labels'].to(device)
15            outputs = iris_model(inputs)
16
17            _, preds = torch.max(outputs, dim=1)
18            loss = criterion(outputs, labels)
19            acc = accuracy(outputs, labels)
20
21            epoch_loss += loss.item()
22            epoch_acc += acc.item()
23
24            # Collect predictions and true labels for metric calculation
25            all_preds.extend(preds.cpu().numpy())
26            all_labels.extend(labels.cpu().numpy())
27
28        # Calculate precision, recall, and F1-score
29        precision = precision_score(all_labels, all_preds, average='weighted')
30        recall = recall_score(all_labels, all_preds, average='weighted')
31        f1 = f1_score(all_labels, all_preds, average='weighted')
32
33        return epoch_loss / len(data_loader), epoch_acc / len(data_loader), precision, recall, f1
34
35 # After training, evaluate the new model on the test set
36 test_loss, test_acc, precision, recall, f1 = test_evaluate(winemodel, test_dataloader, criterion)
37
38 print(f'| Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}% | Precision: {precision:.2f} | Recall: {recall:.2f} | F1-score: {f1:.2f} |')
```

◀ ▶

| Test Loss: 1.435 | Test Acc: 61.25% | Precision: 0.56 | Recall: 0.61 | F1-score: 0.58 |

Summary: by increasing the epoch level to 150 and adding additional layers including increasing neurons. This permitted an increase from an original accuracy of | Test Loss: 1.441 | Test Acc: 60.42% | to achieving an overall result of | Test Loss: 1.459 | Test Acc: 85.94% | which as a sizeable increase of 25.52% accuracy. neruons were increased 100, 200, 200, 150 and 150 to 100. The dramatic increase in neurons and layers allowed for a better accuracy overall.

In []: ▶

```
1 # Stopping Point    Edward Cruz KML614
```

You can view the actual code in my Github: [https://github.com/ecruz0369/DA6233ECruz--2023/blob/main/IS6733kml614_WineLab1\(actual\).ipynb](https://github.com/ecruz0369/DA6233ECruz--2023/blob/main/IS6733kml614_WineLab1(actual).ipynb)